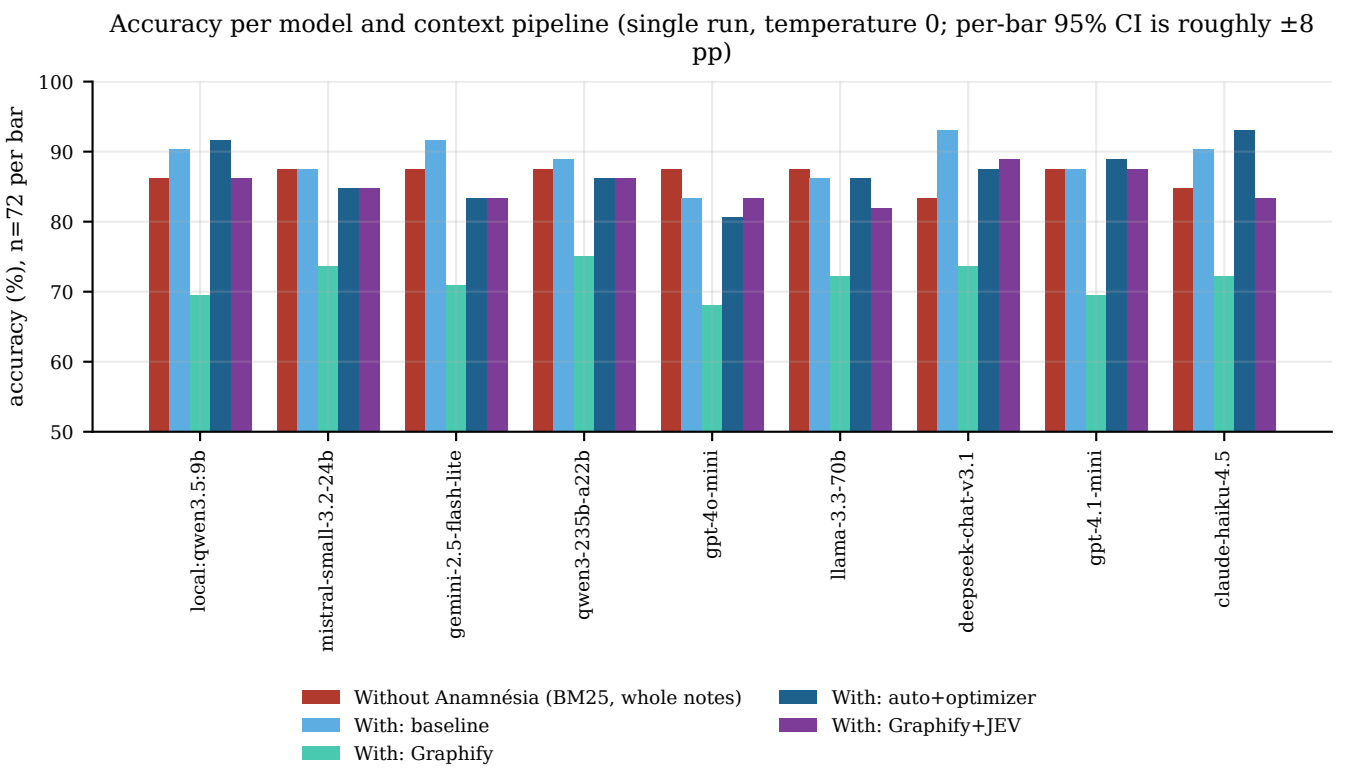
Accuracy per model and context pipeline (single run, temperature 0; per-bar 95% CI is roughly ±8 pp)
accuracy (%), n=72 per bar
100
90
80
70
60
50
local:qwen3.5:9b
mistral-small-3.2-24b
gemini-2.5-flash-lite
qwen3-235b-a22b
gpt-4o-mini
llama-3.3-70b
deepseek-chat-v3.1
gpt-4.1-mini
claude-haiku-4.5
Without Anamnésia (BM25, whole notes)
With: baseline
With: Graphify
With: auto+optimizer
With: Graphify+JEV